

# Homework 3

Jacob Adams

## Table of contents

Question 1	3
Question 2	6
Question 3	9
Question 4	18
<b>Appendix</b>	<b>24</b>

### ! Important

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to only upload a PDF file on Canvas
2. Don't collapse any code cells before submitting.
3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

For this assignment, we will be using the [Wine Quality](#) dataset from the UCI Machine Learning Repository. The dataset consists of red and white *vinho verde* wine samples, from the north of Portugal. The goal is to model wine quality based on physicochemical tests

We will be using the following libraries:

```
library(readr)
library(tidyr)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(purrr)
library(car)
```

Warning: package 'car' was built under R version 4.3.2

Loading required package: carData

Warning: package 'carData' was built under R version 4.3.2

Attaching package: 'car'

The following object is masked from 'package:purrr':

some

The following object is masked from 'package:dplyr':

recode

```
library(glmnet)
```

Warning: package 'glmnet' was built under R version 4.3.2

Loading required package: Matrix

Warning: package 'Matrix' was built under R version 4.3.2

Attaching package: 'Matrix'


The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8

---

## Question 1

 50 points

Regression with categorical covariate and *t*-Test

1.1 (5 points)

Read the wine quality datasets from the specified URLs and store them in data frames `df1` and `df2`.

```
url1 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.csv"
url2 <- "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv"

df1 <- read.csv(url1)
df2 <- read.csv(url2)
```

---

1.2 (5 points)

Perform the following tasks to prepare the data frame `df` for analysis:

1. Combine the two data frames into a single data frame `df`, adding a new column called `type` to indicate whether each row corresponds to white or red wine.
2. Rename the columns of `df` to replace spaces with underscores
3. Remove the columns `fixed_acidity` and `free_sulfur_dioxide`
4. Convert the `type` column to a factor
5. Remove rows (if any) with missing values.

```
df1 <- df1 %>%
  mutate(type = "white")
df2 <- df2 %>%
  mutate(type = "red")
df <- bind_rows(df1, df2)
df <- df %>%
  rename(fixed_acidity = fixed.acidity, volatile_acidity = volatile.acidity, citric_acid = citric.acid)
df <- df %>%
  mutate(fixed_acidity = NULL, free_sulfur_dioxide = NULL) %>%
  mutate(type = as.factor(type))
df <- na.omit(df)
```

Your output to R `dim(df)` should be

```
dim(df)
```

```
[1] 6497    11
```

```
[1] 6497    11
```

---

### 1.3 (20 points)

Recall from STAT 200, the method to compute the  $t$  statistic for the the difference in means (with the equal variance assumption)

1. Using `df` compute the mean of `quality` for red and white wine separately, and then store the difference in means as a variable called `diff_mean`.
2. Compute the pooled sample variance and store the value as a variable called `sp_squared`.
3. Using `sp_squared` and `diff_mean`, compute the  $t$  Statistic, and store its value in a variable called `t1`.

```
df_stats <- df %>%
  group_by(type) %>%
  summarise(mean = mean(quality), sd = sd(quality), n = length(quality))

diff_mean <- df_stats$mean %>%
  diff()
```

```
sp <- sqrt(sum(df_stats$sd^2 * (df_stats$n-1)) / (sum(df_stats$n - 2)) * (1/nrow(df1) + 1/nrow(df2)))
t1 <- diff_mean / sp
t1
```

```
[1] 9.684158
```

---

1.4 (10 points)

Equivalently, R has a function called `t.test()` which enables you to perform a two-sample *t*-Test without having to compute the pooled variance and difference in means.

Perform a two-sample *t*-test to compare the quality of white and red wines using the `t.test()` function with the setting `var.equal=TRUE`. Store the *t*-statistic in `t2`.

```
t_test <- t.test(quality ~ type, data = df, var.equal = TRUE)
t2 <- t_test$statistic
abs(t2)
```

```
      t
9.68565
```

---

1.5 (5 points)

Fit a linear regression model to predict `quality` from `type` using the `lm()` function, and extract the *t*-statistic for the `type` coefficient from the model summary. Store this *t*-statistic in `t3`.

```
fit <- lm(quality ~ type, data = df)
t3 <- coef(summary(fit))[, "t value"][2]
t3
```

```
typewhite
9.68565
```

---

1.6 (5 points)

Print a vector containing the values of `t1`, `t2`, and `t3`. What can you conclude from this? Why?


*From these values we can conclude the t-statistic is very significant*

```
c(t1, t2, t3)
```

```
          t typewhite  
9.684158 -9.685650  9.685650
```

---

## Question 2

 25 points

Collinearity

---

2.1 (5 points)

Fit a linear regression model with all predictors against the response variable `quality`. Use the `broom::tidy()` function to print a summary of the fitted model. What can we conclude from the model summary?

*Based off the p-values of each of the predictors, it seems almost all of them are significant predictors in their own regard. Thus, wine quality is based off a multitude of predictors that can all accurately suggest the quality of wine.*

```
model <- lm(quality ~ ., data = df)  
print(broom::tidy(model))
```

```
# A tibble: 11 x 5
```

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1	(Intercept)	57.5	9.33	6.17	7.44e-10
2	volatile_acidity	-1.61	0.0806	-20.0	4.07e-86
3	citric_acid	0.0272	0.0783	0.347	7.28e- 1
4	residual_sugar	0.0451	0.00416	10.8	3.64e-27
5	chlorides	-0.964	0.333	-2.90	3.78e- 3
6	total_sulfur_dioxide	-0.000329	0.000262	-1.25	2.10e- 1
7	density	-55.2	9.32	-5.92	3.34e- 9
8	pH	0.188	0.0661	2.85	4.38e- 3
9	sulphates	0.662	0.0758	8.73	3.21e-18
10	alcohol	0.277	0.0142	19.5	1.87e-82
11	typewhite	-0.386	0.0549	-7.02	2.39e-12

---

2.2 (10 points)

Fit two **simple** linear regression models using `lm()`: one with only `citric_acid` as the predictor, and another with only `total_sulfur_dioxide` as the predictor. In both models, use `quality` as the response variable. How does your model summary compare to the summary from the previous question?

```
model_citric <- lm(quality ~ citric_acid, data = df)
```

```
model_sulfur <- lm(quality ~ total_sulfur_dioxide, data = df)
```

---

2.3 (5 points)


Visualize the correlation matrix of all numeric columns in `df` using `corrplot()`

```
library(corrplot)
```

Warning: package 'corrplot' was built under R version 4.3.2

corrplot 0.92 loaded

```
num_df <- df[apply(df, is.numeric)]
correlationmatrix <- cor(num_df)
corrplot(correlationmatrix,method = "color")
```



2.4 (5 points)

Compute the variance inflation factor (VIF) for each predictor in the full model using `vif()` function. What can we conclude from this?

*From the VIF of the predictors for the model, we can conclude that most of the predictors have some correlation between each other. The strongest case of collinearity is density, but it is not above 10.*

```
vif(model)
```

volatile_acidity	citric_acid	residual_sugar
2.103853	1.549248	4.680035
chlorides	total_sulfur_dioxide	density
1.625065	2.628534	9.339357
pH	sulphates	alcohol



1.352005	1.522809	3.419849
type		
6.694679		

---

### Question 3

💡 40 points

Variable selection

---

3.1 (5 points)

Run a backward stepwise regression using a `full_model` object as the starting model. Store the final formula in an object called `backward_formula` using the built-in `formula()` function in R

```
backwards_model <- step(model, scope = formula(model), direction = "backward")
```

Start: AIC=-3953.43

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
```

	Df	Sum of Sq	RSS	AIC
- citric_acid	1	0.066	3523.6	-3955.3
- total_sulfur_dioxide	1	0.854	3524.4	-3953.9
<none>			3523.5	-3953.4
- pH	1	4.413	3527.9	-3947.3
- chlorides	1	4.559	3528.1	-3947.0
- density	1	19.054	3542.6	-3920.4
- type	1	26.794	3550.3	-3906.2
- sulphates	1	41.399	3564.9	-3879.5
- residual_sugar	1	63.881	3587.4	-3838.7
- alcohol	1	206.860	3730.4	-3584.8
- volatile_acidity	1	216.549	3740.0	-3567.9

Step: AIC=-3955.3

```
quality ~ volatile_acidity + residual_sugar + chlorides + total_sulfur_dioxide +  
density + pH + sulphates + alcohol + type
```

	Df	Sum of Sq	RSS	AIC
- total_sulfur_dioxide	1	0.818	3524.4	-3955.8
<none>			3523.6	-3955.3
- chlorides	1	4.495	3528.1	-3949.0
- pH	1	4.536	3528.1	-3948.9
- density	1	20.794	3544.4	-3919.1
- type	1	26.943	3550.5	-3907.8
- sulphates	1	41.491	3565.1	-3881.2
- residual_sugar	1	67.371	3590.9	-3834.3
- alcohol	1	235.151	3758.7	-3537.6
- volatile_acidity	1	252.565	3776.1	-3507.5

Step: AIC=-3955.8

```
quality ~ volatile_acidity + residual_sugar + chlorides + density +  
pH + sulphates + alcohol + type
```

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
- pH	1	4.295	3528.7	-3949.9
- chlorides	1	4.523	3528.9	-3949.5
- density	1	21.540	3545.9	-3918.2
- sulphates	1	40.711	3565.1	-3883.2
- type	1	43.664	3568.0	-3877.8
- residual_sugar	1	66.572	3591.0	-3836.2
- alcohol	1	244.545	3768.9	-3521.9
- volatile_acidity	1	256.695	3781.1	-3501.0

```
backward_formula <- formula(backwards_model)
```

---

### 3.2 (5 points)

Run a forward stepwise regression using a `null_model` object as the starting model. Store the final formula in an object called `forward_formula` using the built-in `formula()` function in R

```

null_model <- lm(quality ~ 1, df)
forward_model <- step(null_model, scope = formula(model), direction = "forward")

```

Start: AIC=-1760.04  
quality ~ 1

	Df	Sum of Sq	RSS	AIC
+ alcohol	1	977.95	3975.7	-3186.9
+ density	1	463.41	4490.3	-2396.2
+ volatile_acidity	1	349.71	4604.0	-2233.7
+ chlorides	1	199.47	4754.2	-2025.1
+ type	1	70.53	4883.2	-1851.2
+ citric_acid	1	36.24	4917.4	-1805.7
+ total_sulfur_dioxide	1	8.48	4945.2	-1769.2
+ sulphates	1	7.34	4946.3	-1767.7
+ residual_sugar	1	6.77	4946.9	-1766.9
+ pH	1	1.88	4951.8	-1760.5
<none>			4953.7	-1760.0

Step: AIC=-3186.88  
quality ~ alcohol

	Df	Sum of Sq	RSS	AIC
+ volatile_acidity	1	307.508	3668.2	-3707.9
+ residual_sugar	1	85.662	3890.1	-3326.4
+ type	1	54.335	3921.4	-3274.3
+ citric_acid	1	40.303	3935.4	-3251.1
+ chlorides	1	39.696	3936.0	-3250.1
+ total_sulfur_dioxide	1	31.346	3944.4	-3236.3
+ sulphates	1	7.859	3967.9	-3197.7
+ pH	1	5.938	3969.8	-3194.6
<none>			3975.7	-3186.9
+ density	1	0.005	3975.7	-3184.9

Step: AIC=-3707.89  
quality ~ alcohol + volatile\_acidity

	Df	Sum of Sq	RSS	AIC
+ sulphates	1	48.259	3620.0	-3791.9
+ density	1	38.704	3629.5	-3774.8
+ residual_sugar	1	29.751	3638.5	-3758.8
+ type	1	28.895	3639.3	-3757.3

+ total_sulfur_dioxide	1	5.619	3662.6	-3715.9
+ pH	1	5.533	3662.7	-3715.7
<none>			3668.2	-3707.9
+ chlorides	1	0.162	3668.1	-3706.2
+ citric_acid	1	0.099	3668.1	-3706.1

Step: AIC=-3791.94

quality ~ alcohol + volatile\_acidity + sulphates

	Df	Sum of Sq	RSS	AIC
+ residual_sugar	1	43.989	3576.0	-3869.4
+ density	1	18.661	3601.3	-3823.5
+ type	1	6.012	3614.0	-3800.7
+ chlorides	1	4.988	3615.0	-3798.9
+ citric_acid	1	2.031	3617.9	-3793.6
+ pH	1	1.903	3618.1	-3793.4
<none>			3620.0	-3791.9
+ total_sulfur_dioxide	1	0.817	3619.2	-3791.4

Step: AIC=-3869.37

quality ~ alcohol + volatile\_acidity + sulphates + residual\_sugar

	Df	Sum of Sq	RSS	AIC
+ type	1	20.7581	3555.2	-3905.2
+ total_sulfur_dioxide	1	13.3542	3562.6	-3891.7
+ pH	1	6.6430	3569.3	-3879.5
+ citric_acid	1	4.3384	3571.6	-3875.3
+ chlorides	1	1.8907	3574.1	-3870.8
<none>			3576.0	-3869.4
+ density	1	0.0071	3576.0	-3867.4

Step: AIC=-3905.19

quality ~ alcohol + volatile\_acidity + sulphates + residual\_sugar +  
type

	Df	Sum of Sq	RSS	AIC
+ density	1	20.4623	3534.8	-3940.7
+ chlorides	1	6.6602	3548.6	-3915.4
+ citric_acid	1	5.2242	3550.0	-3912.7
+ pH	1	3.9477	3551.3	-3910.4
+ total_sulfur_dioxide	1	1.2539	3554.0	-3905.5
<none>			3555.2	-3905.2

Step: AIC=-3940.7

quality ~ alcohol + volatile\_acidity + sulphates + residual\_sugar +  
type + density

	Df	Sum of Sq	RSS	AIC
+ chlorides	1	6.0826	3528.7	-3949.9
+ pH	1	5.8541	3528.9	-3949.5
<none>			3534.8	-3940.7
+ citric_acid	1	0.8471	3533.9	-3940.3
+ total_sulfur_dioxide	1	0.5646	3534.2	-3939.7

Step: AIC=-3949.89

quality ~ alcohol + volatile\_acidity + sulphates + residual\_sugar +  
type + density + chlorides

	Df	Sum of Sq	RSS	AIC
+ pH	1	4.2945	3524.4	-3955.8
<none>			3528.7	-3949.9
+ total_sulfur_dioxide	1	0.5765	3528.1	-3948.9
+ citric_acid	1	0.2338	3528.4	-3948.3

Step: AIC=-3955.8

quality ~ alcohol + volatile\_acidity + sulphates + residual\_sugar +  
type + density + chlorides + pH

	Df	Sum of Sq	RSS	AIC
<none>			3524.4	-3955.8
+ total_sulfur_dioxide	1	0.81762	3523.6	-3955.3
+ citric_acid	1	0.02919	3524.4	-3953.9

forward\_model

Call:

lm(formula = quality ~ alcohol + volatile\_acidity + sulphates +  
residual\_sugar + type + density + chlorides + pH, data = df)

Coefficients:

(Intercept)	alcohol	volatile_acidity	sulphates
57.22518	0.28073	-1.62632	0.65234
residual_sugar	typewhite	density	chlorides
0.04425	-0.41760	-54.87625	-0.95067
pH			

0.17589

```
forward_formula <- formula(forward_model)
```

---

### 3.3 (10 points)

1. Create a y vector that contains the response variable (quality) from the df dataframe.
2. Create a design matrix X for the full\_model object using the make\_model\_matrix() function provided in the Appendix.
3. Then, use the cv.glmnet() function to perform LASSO and Ridge regression with X and y.

```
library(glmnet)
y = c(df$quality)
make_model_matrix <- function(formula){
  X <- model.matrix(formula, df)[, -1]
  cnames <- colnames(X)
  for(i in 1:ncol(X)){
    if(!cnames[i] == "typewhite"){
      X[, i] <- scale(X[, i])
    } else {
      colnames(X)[i] <- "type"
    }
  }
  return(X)
}
model_matrix <- make_model_matrix(formula(model))


lasso_model <- cv.glmnet(model_matrix, y, alpha = 1)
ridge_model <- cv.glmnet(model_matrix, y, alpha = 0)
```

Create side-by-side plots of the ridge and LASSO regression results. Interpret your main findings.

*For ridge regression, we should be seeing the sum of squared error tending towards a flat line, and we do. We can see it actually tends towards .55. This means the value of lambda may be too high towards the right and middle end of the graph, or the model is suffering from over fitting. Towards the left tail of the graph, the line tends towards .55 which is what we want to see for a good model. For lasso regression we can see the minimum error occurs*

when  $\log(\lambda) = -7$ . The graph is relatively flat which is a good thing. This means our sum of residual error is expected and not random. Overall, both models level out at a stable mean-squared error. The only noteworthy concern is the standard error is rather high.

```
par(mfrow=c(1, 2))
plot(lasso_model, main = "Lasso Model")
plot(ridge_model, main = "Ridge Model")
```



3.4 (5 points)

Print the coefficient values for LASSO regression at the `lambda.1se` value? What are the variables selected by LASSO?

Store the variable names with non-zero coefficients in `lasso_vars`, and create a formula object called `lasso_formula` using the `make_formula()` function provided in the Appendix.

```
lasso_model$lambda1se
```

NULL

```

lasso_vars <- coef(lasso_model)
make_formula <- function(x){
  as.formula(
    paste("quality ~ ", paste(x, collapse = " + "))
  )
}
print(lasso_vars)

```

11 x 1 sparse Matrix of class "dgCMatrix"

	s1
(Intercept)	5.81837771
volatile_acidity	-0.19128674
citric_acid	.
residual_sugar	0.03943232
chlorides	.
total_sulfur_dioxide	.
density	.
pH	.
sulphates	0.05379620
alcohol	0.36366674
type	.

```

lasso_formula <- make_formula(rownames(lasso_vars)[-1])
lasso_formula

```

```

quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
<environment: 0x00000223c26b4cf8>

```

---

3.5 (5 points)

Print the coefficient values for ridge regression at the `lambda.1se` value? What are the variables selected here? *The variables selected are volatile\_acidity, citric\_acid, residual\_sugar, chlorides, total\_sulfur\_dioxide, density, pH, sulphates, alcohol, and type.*

Store the variable names with non-zero coefficients in `ridge_vars`, and create a formula object called `ridge_formula` using the `make_formula()` function provided in the Appendix.



```
ridge_model$lambda1se
```

NULL

```
ridge_vars <- coef(ridge_model)
print(ridge_vars)
```

11 x 1 sparse Matrix of class "dgCMatrix"

	s1
(Intercept)	5.88519724
volatile_acidity	-0.18384619
citric_acid	0.01816935
residual_sugar	0.10718760
chlorides	-0.04540275
total_sulfur_dioxide	-0.04032981
density	-0.08770244
pH	0.02540356
sulphates	0.08442069
alcohol	0.28041795
type	-0.08863342

```
ridge_formula <- make_formula(rownames(ridge_vars)[-1])
ridge_formula
```

```
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
<environment: 0x00000223c8b8b998>
```

---

3.6 (10 points)

What is the difference between stepwise selection, LASSO and ridge based on you analyses above?

*When we had stepwise regression, the slope of each predictory variable what very high. This is very common in multiple regression, and it is a classic case of overfitting. The LASSO and ridge regression models were able to minimize these slopes to fit the sum of squared error. Thus, allowing the model to be better suited for addition data instead of the df we were provided with.*

---

## Question 4

💡 70 points

Variable selection

---

4.1 (5 points)

Excluding `quality` from `df` we have 10 possible predictors as the covariates. How many different models can we create using any subset of these 10 covariates as possible predictors? Justify your answer. *In the context of choosing all the different potential predictors without replacement and order doesn't matter, we can choose 10 different combinations out of 10 total potential covariates. Thus, the total possible permutations is 10 choose 10.*

```
total <- sum(choose(10, 0:10))
total
```

```
[1] 1024
```

---

4.2 (20 points)

Store the names of the predictor variables (all columns except `quality`) in an object called `x_vars`.

```
x_vars <- colnames(df %>% select(-quality))
```

Use:

- the `combn()` function (built-in R function) and
- the `make_formula()` (provided in the Appendix)

to **generate all possible linear regression formulas** using the variables in `x_vars`. This is most optimally achieved using the `map()` function from the `purrr` package.

```
formulas <- map(
  1:length(x_vars),
  \(x){
    vars <- combn(x_vars, x, simplify = FALSE)
    map(vars, ~ make_formula(.))
  }
) %>% unlist()

formulas <- formulas[!duplicated(formulas)]
```

If your code is right the following command should return something along the lines of:

```
sample(formulas, 4) %>% as.character()
```

```
[1] "quality ~ volatile_acidity + residual_sugar + density + pH"
[2] "quality ~ volatile_acidity + residual_sugar + total_sulfur_dioxide + pH + alcohol + type"
[3] "quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides + total_sulfur_dioxide + pH + alcohol + type"
[4] "quality ~ volatile_acidity + citric_acid + chlorides + total_sulfur_dioxide + density + pH + alcohol + type"
```

```
# Output:
# [1] "quality ~ volatile_acidity + residual_sugar + density + pH + alcohol"
# [2] "quality ~ citric_acid"
# [3] "quality ~ volatile_acidity + citric_acid + residual_sugar + total_sulfur_dioxide + pH + alcohol + type"
# [4] "quality ~ citric_acid + chlorides + total_sulfur_dioxide + pH + alcohol + type"
```

---

#### 4.3 (10 points)

Use `map()` and `lm()` to fit a linear regression model to each formula in `formulas`, using `df` as the data source. Use `broom::glance()` to extract the model summary statistics, and bind them together into a single tibble of summaries using the `bind_rows()` function from `dplyr`.

```
models <- map(formulas, ~lm(.x, data = df))
summaries <- map(models, broom::glance)
single_table = bind_rows(summaries)
```

#### 4.4 (5 points)

Extract the `adj.r.squared` values from `summaries` and use them to identify the formula with the **highest** adjusted R-squared value.

```
get_adj_r_squared <- function(formula, df){
  model2 <- lm(formula, data = df)
  return(summary(model2)$adj.r.squared)
}

#get list of all the adj.r.squared
adj.r.squared <- sapply(summaries, get_adj_r_squared)
```

Store resulting formula as a variable called `rsq_formula`.

```
#get formula where max matches
rsq_formula <- formulas[which.max(adj.r.squared)]
rsq_formula
```

```
[[1]]
quality ~ volatile_acidity
<environment: 0x00000223c94eeb78>
```

---

#### 4.5 (5 points)

Extract the AIC values from `summaries` and use them to identify the formula with the **lowest** AIC value.

```
get_AIC_values <- function(formula, df){
  model2 <- lm(formula, data = df)
  return(summary(model2)$AIC)
}

AIC <- sapply(summaries, function(summary) summary$AIC)
```

Store resulting formula as a variable called `aic_formula`.

```
index <- which.min(AIC)
aic_formula <- formulas[index]
aic_formula
```

```
[[1]]
quality ~ volatile_acidity + residual_sugar + chlorides + density +
      pH + sulphates + alcohol + type
<environment: 0x00000223ca0fde60>
```

---

4.6 (15 points)

Combine all formulas shortlisted into a single vector called `final_formulas`.

```
null_formula <- formula(null_model)
full_formula <- formula(model)

final_formulas <- c(
  null_formula,
  full_formula,
  backward_formula,
  forward_formula,
  lasso_formula,
  ridge_formula,
  rsq_formula,
  aic_formula
)
final_formulas
```

```
[[1]]
quality ~ 1
```

```
[[2]]
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
      total_sulfur_dioxide + density + pH + sulphates + alcohol +
      type
```

```
[[3]]
quality ~ volatile_acidity + residual_sugar + chlorides + density +
      pH + sulphates + alcohol + type
```

```
[[4]]
quality ~ alcohol + volatile_acidity + sulphates + residual_sugar +
      type + density + chlorides + pH
```

```
[[5]]
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
<environment: 0x00000223c26b4cf8>
```

```
[[6]]
quality ~ volatile_acidity + citric_acid + residual_sugar + chlorides +
  total_sulfur_dioxide + density + pH + sulphates + alcohol +
  type
<environment: 0x00000223c8b8b998>
```

```
[[7]]
quality ~ volatile_acidity
<environment: 0x00000223c94eeb78>
```

```
[[8]]
quality ~ volatile_acidity + residual_sugar + chlorides + density +
  pH + sulphates + alcohol + type
<environment: 0x00000223ca0fde60>
```

- Are `aic_formula` and `rsq_formula` the same? How do they differ from the formulas shortlisted in question 3?

*The `aic_formula` and `rsq_formula`'s were not the same, but they did share some covariates. For example, they both had `volatile_acidity` as its  $X_1$  predictory variable. Compared to the Lasso and Ridge they were not similar. They had far less predictory variables.*

- Which of these is more reliable? Why? *The AIC model will be more reliable. The model with the highest  $R$ -squared just indicates its well-suited for the training data. This indicates nothing towards the test data, and it may be a symptom of over fitting.*
- If we had a dataset with 10,000 columns, which of these methods would you consider for your analyses? Why? *I would've probably used lasso or ridge regression in this context, because its automates ways to get rid of unnecessary variables.*

---

#### 4.7 (10 points)

Use `map()` and `glance()` to extract the `sigma`, `adj.r.squared`, `AIC`, `df`, and `p.value` statistics for each model obtained from `final_formulas`. Bind them together into a single data frame `summary_table`. Summarize your main findings.

```
summary_table <- map(
  final_formulas,
  \(x) broom::glance(lm(x, data = df)) %>%
    select(sigma, adj.r.squared, AIC, df, p.value)
) %>% bind_rows()

summary_table %>% knitr::kable()
```

sigma	adj.r.squared	AIC	df	p.value
0.8732553	0.0000000	16679.64	NA	NA
0.7370527	0.2876152	14486.26	10	0
0.7370314	0.2876563	14483.89	8	0
0.7370314	0.2876563	14483.89	8	0
0.7370527	0.2876152	14486.26	10	0
0.7370527	0.2876152	14486.26	10	0
0.8419317	0.0704531	16205.99	1	0
0.7370314	0.2876563	14483.89	8	0

---

## Appendix

### Convenience function for creating a formula object

The following function which takes as input a vector of column names **x** and outputs a **formula** object with **quality** as the response variable and the columns of **x** as the covariates.

```
make_formula <- function(x){
  as.formula(
    paste("quality ~ ", paste(x, collapse = " + "))
  )
}

# For example the following code will
# result in a formula object
# "quality ~ a + b + c"
make_formula(c("a", "b", "c"))
```

### Convenience function for glmnet

The **make\_model\_matrix** function below takes a **formula** as input and outputs a **rescaled** model matrix **X** in a format amenable for **glmnet()**

```
make_model_matrix <- function(formula){
  X <- model.matrix(formula, df)[, -1]
  cnames <- colnames(X)
  for(i in 1:ncol(X)){
    if(!cnames[i] == "typewhite"){
      X[, i] <- scale(X[, i])
    } else {
      colnames(X)[i] <- "type"
    }
  }
  return(X)
}
```



## Session Information

Print your R session information using the following command

```
sessionInfo()
```

```
R version 4.3.1 (2023-06-16 ucrt)
```

```
Platform: x86_64-w64-mingw32/x64 (64-bit)
```

```
Running under: Windows 11 x64 (build 22621)
```

```
Matrix products: default
```

```
locale:
```

```
[1] LC_COLLATE=English_United States.utf8
```

```
[2] LC_CTYPE=English_United States.utf8
```

```
[3] LC_MONETARY=English_United States.utf8
```

```
[4] LC_NUMERIC=C
```

```
[5] LC_TIME=English_United States.utf8
```

```
time zone: America/New_York
```

```
tzcode source: internal
```

```
attached base packages:
```

```
[1] stats      graphics  grDevices  utils      datasets  methods    base
```

```
other attached packages:
```

```
[1] corrplot_0.92 glmnet_4.1-8 Matrix_1.6-5 car_3.1-2 carData_3.0-5
```

```
[6] purrr_1.0.2 dplyr_1.1.2 tidyr_1.3.0 readr_2.1.4
```

```
loaded via a namespace (and not attached):
```

```
[1] jsonlite_1.8.7 compiler_4.3.1 Rcpp_1.0.11 tidyselect_1.2.0
```

```
[5] splines_4.3.1 yaml_2.3.7 fastmap_1.1.1 lattice_0.21-8
```

```
[9] R6_2.5.1 generics_0.1.3 shape_1.4.6 knitr_1.43
```

```
[13] backports_1.4.1 iterators_1.0.14 tibble_3.2.1 pillar_1.9.0
```

```
[17] tzdb_0.4.0 rlang_1.1.1 utf8_1.2.3 broom_1.0.5
```

```
[21] xfun_0.40 cli_3.6.1 withr_2.5.0 magrittr_2.0.3
```

```
[25] digest_0.6.33 foreach_1.5.2 grid_4.3.1 rstudioapi_0.15.0
```

```
[29] hms_1.1.3 lifecycle_1.0.3 vctrs_0.6.3 evaluate_0.21
```

```
[33] glue_1.6.2 codetools_0.2-19 survival_3.5-5 abind_1.4-5
```

```
[37] fansi_1.0.4 rmarkdown_2.24 tools_4.3.1 pkgconfig_2.0.3
```

```
[41] htmltools_0.5.6
```